

Statistics with Recitation: TA Session

Danny Po-Hsien Kang (康柏賢)

November 18, 2025

Today's agenda

1 Data Wrangling

- `summarize()`
- `group_by()`

2 Hypothesis Testing: Chi-Squared Test

- `prop.table()`
- `chisq.test()`

Today's Dataset

- Please download the two datasets from the OpenIntro website.
 - `earthquake.csv`: Record some notable earthquakes that happened in the 20th century.
 - `offshore_drilling.csv`: A survey data randomly asking registered voters in California for their position on drilling for oil and natural gas off the Coast of California.
- After that, import the data

```
earthquake <- read.csv("data/earthquakes.csv")
oil <- read.csv("data/offshore_drilling.csv")
```

Summarize Data: summarize()

- **Data:**

```
> earthquake
```

	year	month	day	richter	area	region	deaths
1	1902	April	19	7.5	...	Guatemala	2000
2	1902	December	16	6.4	Uzbekistan	Russia	4700
3	1903	April	28	7.0	Malazgirt	Turkey	3500
4	1903	May	28	5.8	Gole	Turkey	1000
5	1905	April	4	7.5	Kangra	India	19000

Summarize Data: summarize()

- **Syntax:**

```
dataset %>% summarize(  
  mean_var1 = mean(var1),  
  sd_var1 = sd(var1),  
  max_var2 = max(var2),  
  min_var3 = min(var3),  
  ...  
)
```

Summarize Data: summarize()

- Example:

```
earthquake %>%
  na.omit() %>%
  summarize(
    avg_scale = mean(richter),
    avg_death = mean(deaths)
  )
```

- Output:

	avg_scale	avg_death
1	7.166379	18029.1

Group Data by Variables: group_by()

- **Syntax:**

```
data %>%
  group_by(var) %>%
  summarize(stats = func())
```

- **Example:**

```
earthquake %>%
  group_by(region) %>%
  summarize(
    max_scale = max(richter),
    min_scale = min(richter),
    avg_death = mean(deaths)
  ) %>%
  arrange(desc(avg_death))
```

Group Data by Variables: group_by()

- **Output:**

	region	max_scale	min_scale	avg_death
	<chr>	<dbl>	<dbl>	<dbl>
1	Peru	7.9	7.3	350700
2	Turkmenistan	7.3	7.3	110000
3	China	8	5.9	40236.
4	Armenia	6.8	6.8	25000
5	Italy	7.2	6.5	21950.
6	Japan	8.4	6.9	20324.
7	Pakistan	8	6.2	13100

Group Data by Variables: group_by()

- group_by() adds a grouping attribute to the data.
- Many dplyr verbs then behave per group, not on the full data:
 - summarize(), mutate(), filter(), etc.
- Suppose you want to add summary statistics to the dataframe by region for each row:

```
earthquake %>%  
  group_by(region) %>%  
  mutate(  
    max_scale = max(richter),  
    min_scale = min(richter),  
    avg_death = mean(deaths)  
  ) %>%  
  arrange(desc(avg_death)) %>%  
  ungroup()
```

Group Data by Variables: group_by()

- Output:

	year	deaths	region	max_scale	min_scale	avg_death
	<int>	...	<int>	<chr>	<dbl>	<dbl>
1	1946	...	1400	Peru	7.9	7.3
2	1970	...	700000	Peru	7.9	7.3
3	1948	...	110000	Turk...	7.3	7.3
4	1917	...	1800	China	8	5.9
5	1918	...	1000	China	8	5.9
6	1920	...	200000	China	8	5.9
7	1923	...	3500	China	8	5.9

Group Data by Variables: group_by()

- If you do not want the group structure in data, use ungroup().
 - Forget using ungroup() sometimes can cause tricky problems.
- Suppose you want to find the group with the largest revenue:

```
sales <- data.frame(region = c("N", "N", "S", "S", "S", "S"),
                      store   = c("A", "B", "A", "A", "B", "B"),
                      sales   = c(100, 120, 90, 110, 150, 140))

# forget to ungroup()
by_region_store <- sales %>%
  group_by(region, store) %>%
  summarise(total_sales = sum(sales))

#   region store total_sales
# 1 N       A      100
# 2 N       B      120
# 3 S       A      200
# 4 S       B      290
```

Group Data by Variables: group_by()

- Then you use filter() to select the rows with largest mean:

```
by_region_store %>% filter(total_sales == max(total_sales))
```

- Output:

	region	store	total_sales
1	N	B	120
2	S	B	290

- We get two columns here!
- summarise() defaults to .groups = "drop_last", which drops the last group.
 - by_region_store is now group by region.
- To avoid this, add ungroup() after summarise()
 - or use summarise(..., .groups = "drop")

Create Proportion Table: prop.table()

- Let's first take a look at the oil data:

```
> oil
      v1          v2
1  position college_grad
2    support        yes
3    support        yes
4    support        yes
5    support        yes
6    support        yes
7    support        yes
```

- Our colnames are placed in the first row...
 - This often happens in raw data.

Create Proportion Table: prop.table()

- Quick fix:

```
colnames(oil) <- oil[1, ]
oil <- oil[-1, ]
```

- After fixing:

```
> oil
  position college_grad
2   support      yes
3   support      yes
4   support      yes
5   support      yes
6   support      yes
7   support      yes
```

Create Proportion Table: prop.table()

- Then, let's construct a contingency table:

```
count_table <- table(oil$position,  
                      oil$college_grad  
                     )
```

- Output:**

	no	yes
do_not_know	131	104
oppose	126	180
support	132	154

Create Proportion Table: prop.table()

- **Syntax:**

```
pct_table <- prop.table(contingency_table,  
                         margin = ...)
```

- **Example:**

```
# by row  
pct_table_1 <- prop.table(count_table, margin = 1)  
  
# by column  
pct_table_2 <- prop.table(count_table, margin = 2)
```

Create Proportion Table: prop.table()

- **Output:** pct_table_1 (by row)

	no	yes
do_not_know	0.5574468	0.4425532
oppose	0.4117647	0.5882353
support	0.4615385	0.5384615

- **Output:** pct_table_2 (by column)

	no	yes
do_not_know	0.3367609	0.2374429
oppose	0.3239075	0.4109589
support	0.3393316	0.3515982

Chi-Squared Test: chisq.test()

- **Syntax:**

```
chisq.test(contingency_table)
```

- **Example:**

```
chisq.test(count_table)
```

Chi-Squared Test: chisq.test()

- **Output:**

```
Pearson's Chi-squared test
```

```
data: count_table  
X-squared = 11.461, df = 2, p-value = 0.003246
```

- **Remark:**

- The results of chisq.test() will be saved in a list.
- Don't use the proportion table to perform the Chi-squared test.
- Because the Chi-squared test needs the original count in the data to compute the Chi-squared test statistic.